

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import datasets
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_score
from sklearn.metrics import classification_report
from sklearn.metrics import f1_score
from sklearn.metrics import recall_score
from sklearn.metrics import confusion_matrix
from sklearn.impute import SimpleImputer
from sklearn import tree
from sklearn.decomposition import PCA
from sklearn import metrics
import graphviz
```

## Problem1

```
In [2]: datap1=datasets.load_iris()
```

```
In [3]: type(datap1)
```

```
Out[3]: sklearn.utils.Bunch
```

```
In [4]: datap1.feature_names
```

```
Out[4]: ['sepal length (cm)',
'sepal width (cm)',
'petal length (cm)',
'petal width (cm)']
```

```
In [5]: pdDatap1=pd.DataFrame(np.c_[datap1.data,datap1.target],columns=datap1.feature_
names+['class'])
```

```
In [6]: pdDatap1.head()
```

```
Out[6]:
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	class
0	5.1	3.5	1.4	0.2	0.0
1	4.9	3.0	1.4	0.2	0.0
2	4.7	3.2	1.3	0.2	0.0
3	4.6	3.1	1.5	0.2	0.0
4	5.0	3.6	1.4	0.2	0.0

```
In [7]: pdDatap1.isna().values.any()
```

```
Out[7]: False
```

```
In [8]: pdDatap1.iloc[:, -1].value_counts()
```

```
Out[8]: 2.0    50  
        1.0    50  
        0.0    50  
        Name: class, dtype: int64
```

```

In [9]: xTrainp1,xTestp1,yTrainp1,yTestp1=train_test_split(pdDatap1.iloc[:, :-1],pdData
p1.iloc[:, -1],test_size=0.3)
for i in range(1,6):
    print("For Depth {0} below are the values \n".format(i))
    dtp1=DecisionTreeClassifier(min_samples_split=5,min_samples_leaf=2,max_dep
th=i)
    dtp1.fit(xTrainp1, yTrainp1)
    predictp1=dtp1.predict(xTestp1)
    print(confusion_matrix(yTestp1,predictp1))
    print(classification_report(y_true=yTestp1,y_pred=predictp1,target_names=d
atap1.target_names))
    print('##### \n')
    # print('Micro precision is {0}'.format(precision_score(y_true=yTest,y_pre
d=predict,average='micro')))
    # print('Macro precision is {0}'.format(precision_score(y_true=yTest,y_pre
d=predict,average='macro')))
    # print('Weighted precision is {0}'.format(precision_score(y_true=yTest,y_
pred=predict,average='weighted')))
    # print('##### \n')
    # print('Micro Recall is {0}'.format(recall_score(y_true=yTest,y_pred=pred
ict,average='micro')))
    # print('Macro Recall is {0}'.format(recall_score(y_true=yTest,y_pred=pred
ict,average='macro')))
    # print('Weighted Recal is {0}'.format(recall_score(y_true=yTest,y_pred=pr
edict,average='weighted')))
    # print('##### \n')
    # print('Micro f1Score is {0}'.format(f1_score(y_true=yTest,y_pred=predic
t,average='micro')))
    # print('Macro f1Score is {0}'.format(f1_score(y_true=yTest,y_pred=predic
t,average='macro')))
    # print('Weighted f1Score is {0}'.format(f1_score(y_true=yTest,y_pred=pred
ict,average='weighted')))

```

For Depth 1 below are the values

```
[[15  0  0]
 [ 0 13  0]
 [ 0 17  0]]
precision    recall  f1-score   support

   setosa      1.00      1.00      1.00        15
  versicolor  0.43      1.00      0.60        13
   virginica  0.00      0.00      0.00        17

 accuracy      0.62        45
  macro avg      0.48      0.67      0.53        45
 weighted avg      0.46      0.62      0.51        45

#####
```

For Depth 2 below are the values

```
[[15  0  0]
 [ 0 12  1]
 [ 0  3 14]]
precision    recall  f1-score   support

   setosa      1.00      1.00      1.00        15
  versicolor  0.80      0.92      0.86        13
   virginica  0.93      0.82      0.87        17

 accuracy      0.91        45
  macro avg      0.91      0.92      0.91        45
 weighted avg      0.92      0.91      0.91        45

#####
```

For Depth 3 below are the values

```
[[15  0  0]
 [ 0 12  1]
 [ 0  0 17]]
precision    recall  f1-score   support

   setosa      1.00      1.00      1.00        15
  versicolor  1.00      0.92      0.96        13
   virginica  0.94      1.00      0.97        17

 accuracy      0.98        45
  macro avg      0.98      0.97      0.98        45
 weighted avg      0.98      0.98      0.98        45

#####
```

For Depth 4 below are the values

```
[[15  0  0]
 [ 0 12  1]
 [ 0  0 17]]
precision    recall  f1-score   support
```

setosa	1.00	1.00	1.00	15
versicolor	1.00	0.92	0.96	13
virginica	0.94	1.00	0.97	17
accuracy			0.98	45
macro avg	0.98	0.97	0.98	45
weighted avg	0.98	0.98	0.98	45

#####

For Depth 5 below are the values

```
[[15  0  0]
 [ 0 12  1]
 [ 0  0 17]]
```

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	15
versicolor	1.00	0.92	0.96	13
virginica	0.94	1.00	0.97	17
accuracy			0.98	45
macro avg	0.98	0.97	0.98	45
weighted avg	0.98	0.98	0.98	45

#####

```
C:\Users\venka\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:
1437: UndefinedMetricWarning: Precision and F-score are ill-defined and being
set to 0.0 in labels with no predicted samples.
'precision', 'predicted', average, warn_for)
```

The recall and precision values increase as the height of the tree increases as there is more IG at each level so the maximum Precision and recall are observed at height 2 to 5.

The **f1-score** increases if both the precision and recall are high this can be observed from the height 2-5.

**Micro** calculates the metrics taking the True Positive, True Negative, False Positive, False negative into consideration.

**Macro** calculates the metrics without taking the class imbalance into consideration.

**Weighted** calculates the metrics taking into consideration of weighted imbalances.

## Problem2

```
In [2]: breastCancerDatap2=pd.read_csv("https://archive.ics.uci.edu/ml/machine-learning-
g-databases/breast-cancer-wisconsin/breast-cancer-wisconsin.data",header=None,
names=['Samplenumber','ClumpThickness','Uniformity
yCellSize','UniformCellShape','MarginAdhesion','SingleCellSize','BareNuclei',
'BlandChromatin','NormalNucl','Mitoses','Class'],na_values=['?'])
```

```
In [3]: breastCancerDatap2.isna().values.any()
```

```
Out[3]: True
```

```
In [4]: breastCancerDatap2.isnull().values.any()
```

```
Out[4]: True
```

```
In [5]: print(breastCancerDatap2.dtypes)
```

```
Samplenumber      int64
ClumpThickness    int64
UniformityCellSize int64
UniformCellShape  int64
MarginAdhesion    int64
SingleCellSize    int64
BareNuclei        float64
BlandChromatin    int64
NormalNucl        int64
Mitoses           int64
Class             int64
dtype: object
```

```
In [6]: breastCancerDatap2['BareNuclei'].value_counts()
```

```
Out[6]: 1.0      402
        10.0     132
         5.0      30
         2.0      30
         3.0      28
         8.0      21
         4.0      19
         9.0       9
         7.0       8
         6.0       4
        Name: BareNuclei, dtype: int64
```

```
In [7]: breastCancerDatap2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 699 entries, 0 to 698
Data columns (total 11 columns):
Samplenumber      699 non-null int64
ClumpThickness    699 non-null int64
UniformityCellSize 699 non-null int64
UniformCellShape  699 non-null int64
MarginAdhesion    699 non-null int64
SingleCellSize    699 non-null int64
BareNuclei        683 non-null float64
BlandChromatin    699 non-null int64
NormalNucl        699 non-null int64
Mitoses           699 non-null int64
Class             699 non-null int64
dtypes: float64(1), int64(10)
memory usage: 60.2 KB
```

```
In [8]: imputepr2=SimpleImputer(strategy='mean')
breastCancerDatap2['BareNuclei']=imputepr2.fit_transform(breastCancerDatap2[['BareNuclei']])
```

```
In [9]: xTrainPr2,xTestpr2,yTrainpr2,yTestpr2=train_test_split(breastCancerDatap2.iloc[:, :-1],breastCancerDatap2.iloc[:, -1],test_size=0.3,)
```

```
In [10]: print(xTrainPr2.shape)
print(xTestpr2.shape)
print(yTrainpr2.shape)
print(yTestpr2.shape)
```

```
(489, 10)
```

```
(210, 10)
```

```
(489,)
```

```
(210,)
```

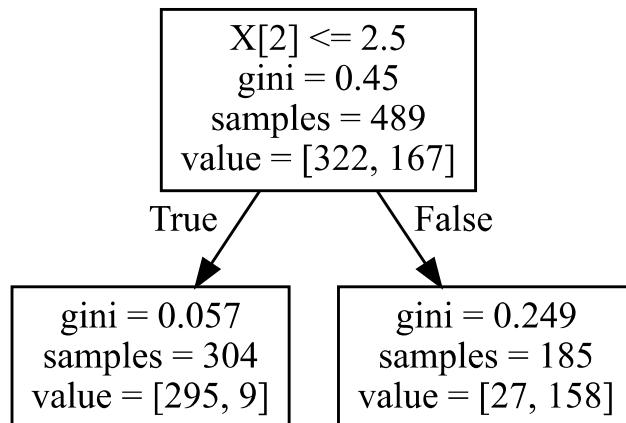
```
In [11]: dtPr2=DecisionTreeClassifier(min_samples_leaf=2,min_samples_split=5,max_depth=1)
dtPr2.fit(xTrainPr2,yTrainpr2)
predictPr2=dtPr2.predict(xTestpr2)
print(classification_report(y_true=yTestpr2,y_pred=predictPr2))
print(confusion_matrix(y_true=yTestpr2,y_pred=predictPr2))
```

	precision	recall	f1-score	support
2	0.98	0.90	0.93	136
4	0.84	0.96	0.89	74
accuracy			0.92	210
macro avg	0.91	0.93	0.91	210
weighted avg	0.93	0.92	0.92	210

```
[[122 14]
 [ 3 71]]
```

```
In [12]: graphviz.Source(tree.export_graphviz(dtPr2))
```

Out[12]:



```
In [13]: dtPr2v2=DecisionTreeClassifier(min_samples_leaf=2,min_samples_split=5,max_dept
h=2)
dtPr2v2.fit(xTrainPr2,yTrainpr2)
predictPr2v2=dtPr2v2.predict(xTestpr2)
print(classification_report(y_true=yTestpr2,y_pred=predictPr2v2))
print(confusion_matrix(y_true=yTestpr2,y_pred=predictPr2v2))
```

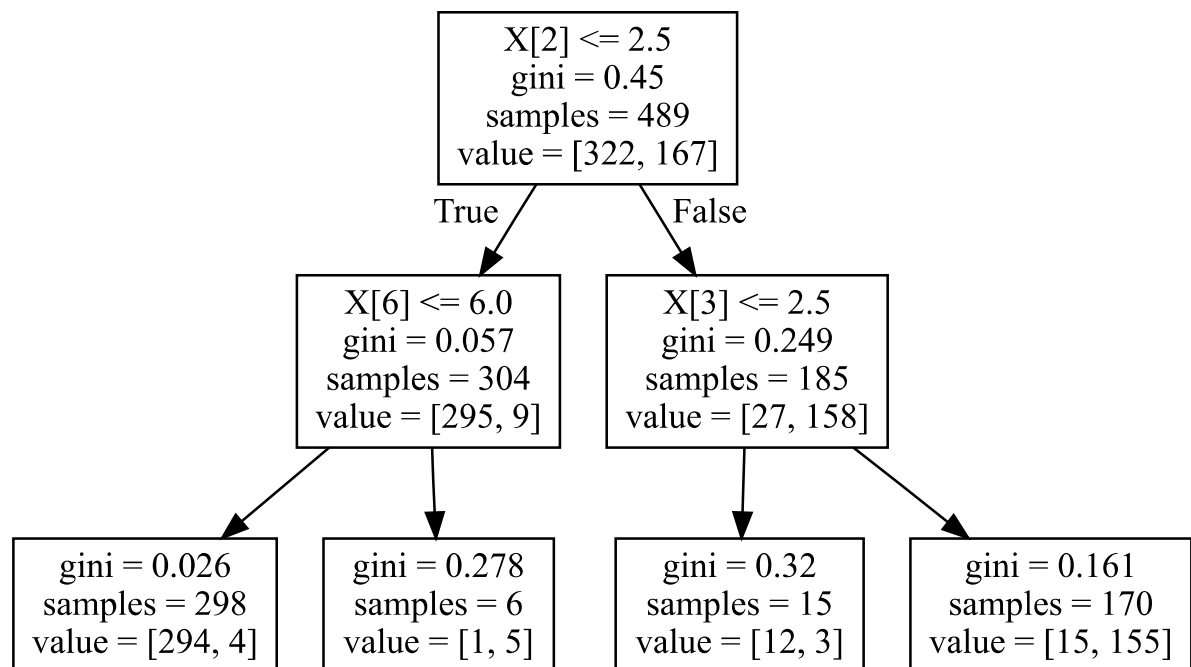
	precision	recall	f1-score	support
2	0.97	0.94	0.96	136
4	0.90	0.95	0.92	74
accuracy			0.94	210
macro avg	0.93	0.94	0.94	210
weighted avg	0.94	0.94	0.94	210

```
[[128  8]
 [ 4 70]]
```

```
In [14]: graphviz.Source(tree.export_graphviz(dtPr2v2))
```

Out[14]:



## Problem 3

```
In [2]: dataPr3=pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-database
s/breast-cancer-wisconsin/wdbc.data',names=['id', 'diagnosis','radiusMean','te
xtureMean','perimeterMean','areaMean','smoothnessMean','compactnessMean', 'con
cavityMean','oncavePointsMean', 'symmetryMean', 'fractalDimensionMean','radiusSe',
'textureSe', 'perimeterSe', 'areaSe', 'smoothnessSe','compactnessSe', 'co
ncavitySe', 'concave pointsSe', 'symmetrySe','fractalDimensionSe', 'radiusWorst',
'textureWorst','perimeterWorst', 'areaWorst', 'smoothnessWorst','compactne
ssWorst', 'concavityWorst', 'concavePointsWorst','symmetryWorst', 'fractalDime
nsionWorst'])
```



In [3]: dataPr3.head()

Out[3]:

	id	diagnosis	radiusMean	textureMean	perimeterMean	areaMean	smoothnessMean
0	842302	M	17.99	10.38	122.80	1001.0	0.11840
1	842517	M	20.57	17.77	132.90	1326.0	0.08474
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960
3	84348301	M	11.42	20.38	77.58	386.1	0.14250
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030

5 rows × 32 columns



In [4]: dataPr3.iloc[:, -1].shape

Out[4]: (569,)

In [5]: dataPr3[["diagnosis"]].shape

Out[5]: (569, 1)

In [6]: pca=PCA(n\_components=1)

In [7]: principalComponents=pca.fit\_transform(dataPr3.iloc[:, 2:])  
 print(pca.explained\_variance\_ratio\_)  
 print(principalComponents.shape)  
 print(principalComponents[:, 0].shape)

[0.98204467]  
 (569, 1)  
 (569,)

In [8]: xTrainPr3, xTestPr3, yTrainPr3, yTestPr3=train\_test\_split(principalComponents, dataPr3[["diagnosis"]], test\_size=0.3)

In [9]: dtPr3=DecisionTreeClassifier(max\_depth=1, min\_samples\_split=5, min\_samples\_leaf=2)  
 dtPr3.fit(xTrainPr3, yTrainPr3)  
 predict=dtPr3.predict(xTestPr3)  
 print(classification\_report(y\_pred=predict, y\_true=yTestPr3))

	precision	recall	f1-score	support
B	0.90	0.97	0.94	114
M	0.94	0.79	0.86	57
accuracy			0.91	171
macro avg	0.92	0.88	0.90	171
weighted avg	0.91	0.91	0.91	171

```
In [10]: print(confusion_matrix(y_pred=predict,y_true=yTestPr3))
```

```
[[111  3]
 [ 12 45]]
```

```
In [11]: pca2=PCA(n_components=1)
principalComponentsv2=pca.fit_transform(dataPr3.iloc[:,2:])
```

```
In [12]: xTrainPr3v2,xTestPr3v2,yTrainPr3v2,yTestPr3v2=train_test_split(principalComponentsv2,dataPr3[["diagnosis"]],test_size=0.3)
```

```
In [13]: dtPr3v2=DecisionTreeClassifier(max_depth=1, min_samples_split=5, min_samples_leaf=2)
dtPr3v2.fit(xTrainPr3v2,yTrainPr3v2)
predictv2=dtPr3v2.predict(xTestPr3v2)
print(classification_report(y_pred=predict,y_true=yTestPr3v2))
```

	precision	recall	f1-score	support
B	0.64	0.71	0.67	112
M	0.31	0.25	0.28	59
accuracy			0.55	171
macro avg	0.48	0.48	0.48	171
weighted avg	0.53	0.55	0.54	171

```
In [14]: print(confusion_matrix(y_pred=predictv2,y_true=yTestPr3v2))
```

```
[[109  3]
 [ 12 47]]
```

## Problem4

```
In [2]: feature1p4=np.random.normal(5,2,500)
feature2p4=np.random.normal(-5,2,500)
print(feature1p4.shape)
print(feature2p4.shape)
class1p4=np.repeat(0,500)
class2p4=np.repeat(1,500)
data1p4=pd.DataFrame({'feature':feature1p4,'classLabel':class1p4})
data2p4=pd.DataFrame({'feature':feature2p4,'classLabel':class2p4})

(500,)
(500,)
```

In [3]: data1p4.head()

Out[3]:

	feature	classLabel
0	6.000080	0
1	0.550831	0
2	3.645289	0
3	6.771784	0
4	5.445052	0

In [4]: data2p4.head()

Out[4]:

	feature	classLabel
0	-8.339510	1
1	-4.475830	1
2	-1.711105	1
3	-5.510317	1
4	-5.583594	1

In [5]: finalDatap4=pd.concat([data1p4,data2p4])  
finalDatap4.head()

	feature	classLabel
0	6.000080	0
1	0.550831	0
2	3.645289	0
3	6.771784	0
4	5.445052	0
(1000, 2)		

In [6]: xTrainP4,xTestP4,yTrainP4,yTestP4=train\_test\_split(finalDatap4[['feature']],finalDatap4['classLabel'],test\_size=0.3)

In [7]: dtp4 = DecisionTreeClassifier(max\_depth=2)  
dtp4.fit(xTrainP4, yTrainP4)  
predictp4=dtp4.predict(xTestP4)

In [14]: fprp4,tprp4,thresholdsp4=metrics.roc\_curve(y\_score=predictp4,y\_true=yTestP4)

In [15]: print("THRESHOLD is {0}".format(thresholdsp4))

THRESHOLD is [2 1 0]

## Recitation

## Chapter 2

### Exercise 2)

1) Count of C0 is 10

count of C1 is 10

$$\text{Gini} = 1 - (10/20)^2 - (10/20)^2 \\ = 0.5$$

### Exercise 2)

2) As ID are unique probability of each id is  $1/20$

As there are only two possibilities of classes for the each ID Gini index of each id is 0

so total Gini Index  $= 20(1/20)0 = 0$

### Exercise 2)

3)

Female count=10

Male count=10

$$\text{Gini} = 1 - (10/20)^2 - (10/20)^2 \\ = 0.5$$

### Exercise 2)

4)

Family car Type Probability=  $4/20 = 0.2$

Sports car Type Probability=  $8/20 = 0.4$

Luxury car Type Probability=  $8/20 = 0.4$

$$P(\text{Family}|C1) = 0.75 \quad P(\text{Family}|C0) = 0.25$$

$$\text{gini} = 1 - (0.75)^2 - (0.25)^2 \\ = 1 - 0.5625 - 0.0625 \\ = 0.375$$

$$P(\text{Sports}|C1) = 0 \quad P(\text{Sports}|C0) = 1$$

$$\text{gini} = 1 - (0)^2 - (1)^2 \\ = 0$$

$$P(\text{Luxury}|C1) = 0.125 \quad P(\text{Luxury}|C0) = 0.875$$

$$\text{gini} = 1 - (0.125)^2 - (0.875)^2 \\ = 1 - 0.015625 - 0.765625 \\ = 0.21875$$

$$\text{Gini Index} = (0.20 \cdot 0.375) + (0.40) + (0.4 \cdot 0.21875) = 0.1625$$

**Exercise 2)****5)**Small Shirt size:  $5/20=0.25$ Medium Shirt size:  $7/20=0.35$ Large Shirt size:  $4/20=0.2$ Extra Large size:  $4/20=0.2$  $P(\text{Small}|C1) = 2/5=0.4$   $P(\text{Small}|C0) = 3/5=0.6$  $Gini = 1 - (0.4)^2 - (0.6)^2$  $= 0.48$  $P(\text{Medium}|C1) = 4/7=0.5714$   $P(\text{Medium}|C0) = 3/7=0.428$  $Gini = 1 - (0.5714)^2 - (0.428)^2$  $= 0.49$  $P(\text{Large}|C1) = 2/4=0.5$   $P(\text{Large}|C0) = 2/4=0.5$  $Gini = 1 - (0.5)^2 - (0.5)^2$  $= 0.5$  $P(\text{ExtraLarge}|C1) = 2/4=0.5$   $P(\text{ExtraLarge}|C0) = 2/4=0.5$  $Gini = 1 - (0.5)^2 - (0.5)^2$  $= 0.5$  $Gini\ Index = (0.25 \cdot 0.48) + (0.35 \cdot 0.49) + (0.2 \cdot 0.5) + (0.2 \cdot 0.5)$  $= 0.4915$ **Exercise 2)****6)**Lower the gini index better it is from above **carType** has lower Gini i.e. **Gini of carType=0.1625****Exercise 2)****7)**

It is just a primary key needed for RDBMS to store data it holds no rela meaning.

**Exercise 6)****1)** $Gini\ index = 1 - (3/10)^2 - (7/10)^2$  $= 0.42$  $Miscalssification\ Error = 1 - (7/10)$  $= 0.3$

**Exercise 6)****2)**

$$\text{ginic1} = 1 - 1 - (3/3)^2 - (0/3)^2$$

$$= 0$$

$$\text{ginic2} = 1 - (4/7)^2 - (3/7)^2$$

$$= 0.489$$

**Exercise 6)****3)**

$$\text{misClassificationC1} = 1 - \text{Max}((3/3), (0/3)) = 0 \quad \text{misClassificationC2} = 1 - \text{Max}((4/7), (3/7)) = 1 - (4/7) = 0.428$$

In [ ]: